Name: Ceren Kurt

Student Number: 243910

Socratic Dialogue Preparation Document

Statement

Should AI developers bear legal responsibility for representational harms caused by their systems, even in the absence of malicious intent?

Phase 1: Thoughts and Opinions

When I first read this statement, I immediately thought: yes, they should be responsible. If AI systems are causing harm to certain groups in society, even if it's unintentional, there must be accountability. Developers are the ones building these systems, and they should be aware of the possible consequences. At the same time, I hesitated. Technology is complicated, and developers can't always predict every possible outcome. It didn't seem entirely fair to punish someone for harm they didn't intend to cause. My initial opinion was conflicted: I believe responsibility is important, but it must also be balanced and realistic.

Phase 2: Internal Dialogue (Research and Informed Position)

To build a more informed position, I began researching examples of representational harm in AI systems. One well-known case is Amazon's AI recruitment tool, which unfairly penalized women because it was trained on biased historical hiring data. Another example is COMPAS, a risk assessment algorithm used in U.S. courts that was found to give higher risk scores to Black defendants compared to white defendants. These examples show how AI can reinforce harmful stereotypes and social inequalities, even without explicit intent from developers. From an ethical standpoint, scholars like Kate Crawford ("Atlas of AI") argue that AI is not neutral-it reflects the values and biases of those who build it. Similarly, Microsoft's Responsible AI Guidelines and IBM's AI Ethics documentation both emphasize that developers have a duty to proactively address bias and test systems for potential harms. These practices include data audits, inclusive testing, and explainability. Based on this evidence, my informed position is: Yes, AI developers should bear some legal responsibility, particularly when they neglect to follow known best practices for ethical AI development. Accountability can act as a motivation to apply more rigorous ethical checks throughout the design process.

Phase 3: Dialogue with the Other (Steelmanning Opposing View)

The strongest counterargument is that AI development is a collaborative process involving large teams and sometimes third-party tools or datasets. It is unreasonable to hold individual developers legally responsible for outcomes they cannot fully control. Furthermore, the technology is still evolving. Mistakes are often unintentional, and punishing developers might discourage innovation and experimentation. Another point is that intent matters. If a developer did not act with malice and took all reasonable steps to ensure fairness, holding them legally accountable could set an unrealistic standard. This argument suggests that instead of assigning legal blame, we should focus on system-level reforms-such as mandatory impact assessments, stronger regulation, and clearer AI governance. I learned from this view that over-punishing developers could create fear and slow progress. Not all harm is the result of negligence; sometimes it is a result of the inherent complexity of AI systems. This perspective helped me see the importance of collective responsibility and better structural frameworks.

Phase 4: Internal Reflection

Reflecting on both sides, I realized that a balanced view is needed. While I still believe developers should be held accountable, I no longer think it should always involve legal punishment. There should be a distinction between negligence and unintended consequences despite good-faith efforts. If a developer or team failed to implement basic fairness testing or ignored known risks, then yes, legal responsibility may be justified. However, if they followed all reasonable protocols and still harm occurred, the responsibility should shift toward improving the overall regulatory framework. This phase made me realize that responsibility in AI should not only be about blaming individuals but also about building a culture of accountability and care in development teams.

Phase 5: Temporary Conclusion

In conclusion, AI developers should bear legal responsibility for representational harms only when those harms result from negligence or failure to follow ethical practices. If harm occurs despite developers making good-faith efforts to follow responsible AI guidelines, then legal punishment is not appropriate. Instead, we should focus on building stronger systems of governance and accountability in AI development. This position balances the need to protect individuals and communities from harm with the need to support ethical innovation. It encourages developers to think critically about their role in shaping AI's social impact, without placing unfair burdens on those who act responsibly.

References

- Crawford, K. (2021). Atlas of AI: Power, Politics, and the Planetary Costs of Artificial Intelligence. Yale University Press.

- Binns, R. (2018). Fairness in Machine Learning: Lessons from Political Philosophy. Communications of the ACM, 61(9), 8-9.

- Raji, I. D., & Buolamwini, J. (2019). Actionable Auditing: Investigating the Impact of Publicly Naming Biased Performance Results of Commercial AI Products. Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society.

- Microsoft. (2021). Responsible AI Standard v2.

- IBM. (2020). Everyday Ethics for Artificial Intelligence.

- O'Neil, C. (2016). Weapons of Math Destruction. Crown Publishing Group